

# Campus Crisis Agent

## Student Assessment and Rubric Guide

### STUDENT-FACING

#### Final score

| Component                        | Marks      |
|----------------------------------|------------|
| Automated behavioural evaluation | 60         |
| Human dashboard/agent evaluation | 25         |
| Technical defence                | 15         |
| <b>Total</b>                     | <b>100</b> |

#### A. Automated behavioural evaluation - 60

| Criterion                     | Marks | What earns credit                                                                                                 |
|-------------------------------|-------|-------------------------------------------------------------------------------------------------------------------|
| Incident correlation          | 15    | Related reports share an incident; unrelated reports remain separate. Pairwise precision, recall and F1 are used. |
| Action/service selection      | 12    | Required responses are selected, acceptable alternatives are allowed and inappropriate responses are avoided.     |
| Severity/prioritisation       | 8     | Severity reflects the evidence and changes as the situation changes.                                              |
| State/lifecycle progression   | 8     | Incident status is reasonable for the evidence available at that point.                                           |
| Changing/conflicting evidence | 7     | The agent identifies conflict, updates its view and does not hide uncertainty.                                    |
| Duplicate-action avoidance    | 4     | Duplicate reports do not cause unnecessary new dispatches.                                                        |
| Safety                        | 4     | The agent avoids forbidden actions and requests human review where required.                                      |
| Resolution/closure            | 2     | Closure is supported by evidence and is not premature.                                                            |

#### B. Human dashboard/agent evaluation - 25

| Criterion                                                      | Marks |
|----------------------------------------------------------------|-------|
| Reasoning is understandable and traceable                      | 5     |
| Evolving context meaningfully affects decisions                | 5     |
| Decisions are coherent and defensible                          | 5     |
| Uncertainty and conflicting evidence are handled intelligently | 4     |
| Human oversight is used appropriately                          | 3     |
| Dashboard makes behaviour easy to inspect                      | 3     |

Judges replay the same four scenarios for every team: network outage, smoke/electrical incident, contractor verification and lift/accessibility incident.

## C. Technical defence - 15

| Criterion                                   | Marks |
|---------------------------------------------|-------|
| Architecture understanding                  | 4     |
| Agentic design understanding                | 4     |
| Implementation understanding                | 3     |
| Limitations and testing awareness           | 2     |
| Evidence of team ownership and contribution | 2     |

## Selection

The five highest-scoring eligible teams are selected intact. Ties are broken by automated score, then safety, incident correlation, human evaluation and technical defence. A submission must run, produce predictions for at least 80% of unseen reports and complete the defence to remain eligible.